



UNCERTAINTY - PROBABILITY



UNCERTAINTY

Let action A_t = leave for airport t minutes before flight
Will A_t get me there on time?

Problems:

1. partial observability (road state, etc.)
2. multi-agent problem (other drivers' plans)
3. noisy sensors (uncertain traffic reports)
4. uncertainty in action outcomes (flat tire, etc.)
5. immense complexity of modeling and predicting traffic

UNCERTAINTY

Let action A_t = leave for airport t minutes before flight Will A_t get me there on time?

Problems:

1. partial observability (road state, etc.)
2. multi-agent problem (other drivers' plans)
3. noisy sensors (uncertain traffic reports)
4. uncertainty in action outcomes (flat tire, etc.)
5. immense complexity of modeling and predicting traffic

Hence a purely logical approach either

1. risks falsehood: “ A_{25} will get me there on time”, or
2. leads to conclusions that are too weak for decision making:

“ A_{25} will get me there on time if there's no accident on the bridge and it doesn't rain and my tires remain intact, etc., etc.”

“ A_{1440} should get me there on time but I'd have to stay overnight in the airport.”

PROPOSITIONAL LOGIC AND PROBABILITY

Their ontological commitments are the same

The world is a set of facts that do or do not hold

Ontology is the philosophical study of the nature of being, becoming, existence, or reality; what exists in the world?

Their epistemological commitments differ

- **Logic agent** believes true, false, or no opinion
- **Probabilistic agent** has a numerical degree of belief between 0 (false) and 1 (true)
- **Epistemology** is the philosophical study of the nature and scope of knowledge; how, and in what way, do we know about the world?

UNCERTAINTY IN THE WORLD

- Uncertainty due to
 - Randomness
 - Overwhelming complexity
 - Lack of knowledge
 - ...

UNCERTAINTY IN THE WORLD

- Uncertainty due to
 - Randomness
 - Overwhelming complexity
 - Lack of knowledge
 - ...
- Probability gives
 - natural way to describe our assumptions
 - rules for how to combine information

UNCERTAINTY IN THE WORLD

- Uncertainty due to
 - Randomness
 - Overwhelming complexity
 - Lack of knowledge
 - ...
- Probability gives
 - natural way to describe our assumptions
 - rules for how to combine information
- Subjective probability
 - Relate to agent's own state of knowledge: $P(A_{25} \mid \text{no accidents}) = 0.05$
 - Not assertions about the world; indicate **degrees of belief**
 - Change with new evidence: $P(A_{25} \mid \text{no accidents, 5am}) = 0.20$

MAKING DECISIONS UNDER UNCERTAINTY

- Suppose I believe the following:
 - $P(A_{25} \text{ gets me there on time} \mid \dots) = 0.04$
 - $P(A_{90} \text{ gets me there on time} \mid \dots) = 0.70$
 - $P(A_{120} \text{ gets me there on time} \mid \dots) = 0.95$
 - $P(A_{1440} \text{ gets me there on time} \mid \dots) = 0.9999$
- Which action to choose?

MAKING DECISIONS UNDER UNCERTAINTY

- Which action to choose?
- Depends on my **preferences** for missing flight vs. time spent waiting, etc.
 - **Utility theory** is used to represent and infer preferences
 - **Decision theory**= probability theory + utility theory

- **Expected utility** of action a in state s

$$= \sum_{\text{outcome in Results}(s,a)} P(\text{outcome}) * \text{Utility}(\text{outcome})$$

- A rational agent acts to maximize expected utility

EX: AIRPORT

- Suppose I believe the following:
 - $P(\text{A25 gets me there on time} \mid \dots) = 0.04$
 - $P(\text{A90 gets me there on time} \mid \dots) = 0.70$
 - $P(\text{A120 gets me there on time} \mid \dots) = 0.95$
 - $P(\text{A1440 gets me there on time} \mid \dots) = 0.9999$
 - $\text{Utility}(\text{on time}) = \$1,000$
 - $\text{Utility}(\text{not on time}) = -\$10,000$

- **Expected utility** of action a in state s

$$= \sum_{\text{outcome in Results}(s,a)} P(\text{outcome}) * \text{Utility}(\text{outcome})$$

$$E(\text{Utility}(\text{A25})) = 0.04 * \$1,000 + 0.96 * (-\$10,000) = -\$9,560$$

$$E(\text{Utility}(\text{A90})) = 0.7 * \$1,000 + 0.3 * (-\$10,000) = -\$2,300 \quad E(\text{Utility}(\text{A120}))$$

$$= 0.95 * \$1,000 + 0.05 * (-\$10,000) = \$450 \quad E(\text{Utility}(\text{A1440})) =$$

$$0.9999 * \$1,000 + 0.0001 * (-\$10,000) = \$998.90$$

- Have not yet accounted for disutility of staying overnight at the airport, etc.

PROBABILITY

- $P(a)$ is the probability of proposition “a”
 - E.g., $P(\text{it will rain in London tomorrow})$
 - The proposition “a” is actually true or false in the real world
 - $P(a)$ is our degree of belief that proposition “a” is true in the real world
 - $P(a)$ = “prior” or marginal or unconditional probability
 - Assumes no other information is available
- **Axioms of probability:**
 - $0 \leq P(a) \leq 1$
 - $P(\text{NOT}(a)) = 1 - P(a)$
 - $P(\text{true}) = 1$
 - $P(\text{false}) = 0$
 - $P(a \text{ OR } b) = P(a) + P(b) - P(a \text{ AND } b)$

INTERPRETATIONS OF PROBABILITY

- **Relative Frequency:** *Usually taught in school*
 - $P(a)$ represents the frequency that event a will happen in repeated trials.
 - Requires event a to have happened enough times for data to be collected.
- **Degree of Belief:** *A more general view of probability*
 - $P(a)$ represents an agent's degree of belief that event a is true.
 - Can predict probabilities of events that occur rarely or have not yet occurred.
 - Does not require new or different rules, just a different interpretation.
- Examples:
 - a = "life exists on another planet"
 - What is $P(a)$? We all will assign different probabilities
 - a = "California will secede from the US"
 - What is $P(a)$?
 - a = "over 50% of the students in this class will get A's"
 - What is $P(a)$?

CONCEPTS OF PROBABILITY

- Unconditional Probability
 - $P(a)$, the probability of “a” being true, or $P(a=\text{True})$
 - Does not depend on anything else to be true (**unconditional**)
 - Represents the probability prior to further information that may adjust it (**prior**)
 - Also sometimes “**marginal**” probability (vs. joint probability)


CONCEPTS OF PROBABILITY

- Conditional Probability
 - $P(a|b)$, the probability of “a” being true, given that “b” is true
 - Relies on “b” = true (**conditional**)
 - Represents the prior probability adjusted based upon new information “b” (**posterior**)
 - Can be generalized to more than 2 random variables:
 - e.g. $P(a|b, c, d)$

CONCEPTS OF PROBABILITY

- Conditional Probability

- $P(a|b)$, the probability of “a” being true, given that “b” is true
- Relies on “b” = true (**conditional**)
- Represents the prior probability adjusted based upon new information “b” (**posterior**)
- Can be generalized to more than 2 random variables:
 - e.g. $P(a|b, c, d)$



We often use comma to abbreviate AND.

- Joint Probability

- $P(a, b) = P(a \wedge b)$, the probability of “a” and “b” both being true
- Can be generalized to more than 2 random variables:
 - e.g. $P(a, b, c, d)$

RANDOM VARIABLES

- **Random Variable:**
 - Basic element of probability assertions
 - Similar to CSP variable, but values reflect probabilities not constraints.
 - Variable: A
 - Domain: $\{a_1, a_2, a_3\}$ <-- events / outcomes
- Types of Random Variables:
 - **Boolean** random variables : $\{ true, false \}$
 - e.g., *Cavity* (= do I have a cavity?)
 - **Discrete** random variables : one value from a set of values
 - e.g., *Weather* is one of $\{sunny, rainy, cloudy, snow\}$
 - **Continuous** random variables : a value from within constraints
 - e.g., *Current temperature* is bounded by $(10^\circ, 200^\circ)$
- Domain values must be **exhaustive and mutually exclusive**:
 - One of the values must always be the case (**Exhaustive**)
 - Two of the values cannot both be the case (**Mutually Exclusive**)

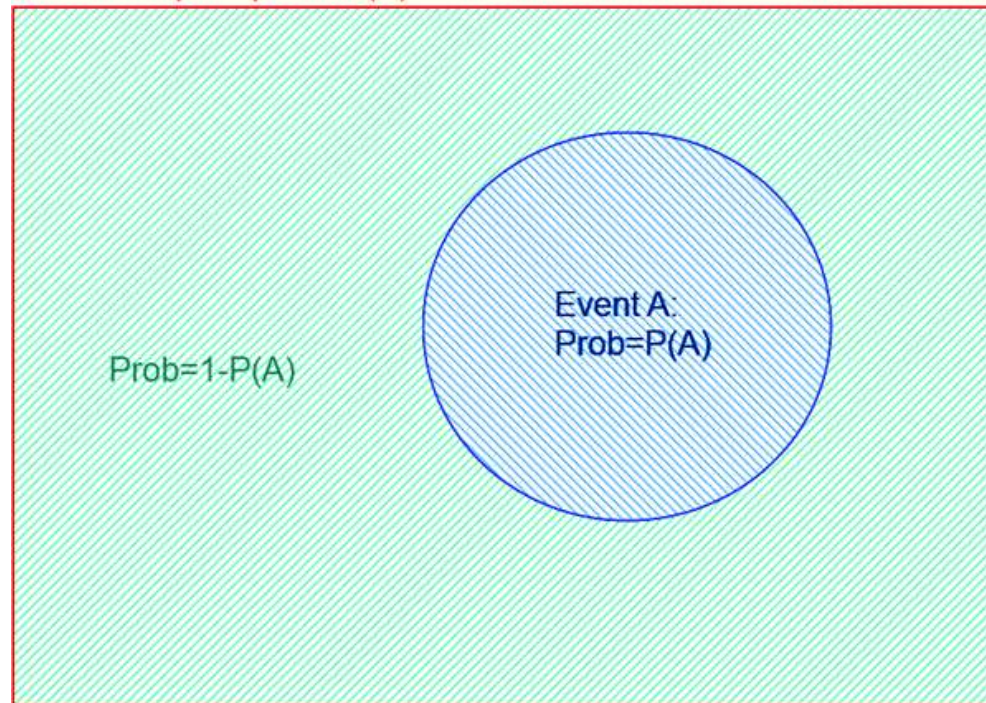
RANDOM VARIABLES

- **Example:** Coin flip
 - Variable = R, the result of the coin flip
 - Domain = {heads, tails, edge} <-- must be exhaustive
 - $P(R = \text{heads}) = 0.4999$
 - $P(R = \text{tails}) = 0.4999$ } <-- must be exclusive
 - $P(R = \text{edge}) = 0.0002$ }
- Shorthand is often used for simplicity:
 - Upper-case letters for variables, lower-case letters for values.
 - E.g., $P(A) \equiv \langle P(A=a_1), P(A=a_2), \dots, P(A=a_n) \rangle$ for all n values in Domain(A)
 - E.g., $P(a) \equiv P(A = a)$
 $P(a|b) \equiv P(A=a | B=b)$
 $P(a, b) \equiv P(A=a \wedge B = b)$
- Two kinds of probability propositions:
 - **Elementary propositions** are an assignment of a value to a random variable:
 - e.g., *Weather = sunny*; e.g., *Cavity = false* (abbreviated as $\neg \text{cavity}$)
 - **Complex propositions** are formed from elementary propositions and standard logical connectives :
 - e.g., *Cavity = false \vee Weather = sunny*

PROBABILITY SPACE

$$P(A) + P(\neg A) = 1$$

Entire Sample Space: $P(S)=1$




Area = Probability of Event

AND PROBABILITY

$$P(A, B) = P(A \wedge B) = P(A) + P(B) - P(A \vee B)$$

Entire Sample Space: $P(S)=1$




Area = Probability of Event

OR PROBABILITY

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$

Entire Sample Space: $P(S)=1$




Area = Probability of Event

CONDITIONAL PROBABILITY

We often use comma to abbreviate AND.

$$P(A | B) = P(A, B) / P(B)$$

Entire Sample Space: $P(S)=1$




Area = Probability of Event

PRODUCT RULE

$$P(A,B) = P(A|B) P(B)$$

Entire Sample Space: $P(S)=1$



Area = Probability of Event

USING THE PRODUCT RULE

- **Applies to any number of variables:**
 - $P(a, b, c) = P(a, b | c) P(c) = P(a | b, c) P(b, c)$
 - $P(a, b, c | d, e) = P(a | b, c, d, e) P(b, c | d, e)$
- **Factoring:** (AKA **Chain Rule** for probabilities)

- By the product rule, we can always write:

$$P(a, b, c, \dots z) = P(a | b, c, \dots z) P(b, c, \dots z)$$

We often use comma to abbreviate AND.

- Repeatedly applying this idea, we can write:

$$P(a, b, c, \dots z) = P(a | b, c, \dots z) P(b | c, \dots z) P(c | \dots z) \dots P(z)$$

- This holds for any ordering of the variables

INFERENCE BY ENUMERATION

Start with the joint distribution:

	<i>toothache</i>		\neg <i>toothache</i>	
	<i>catch</i>	\neg <i>catch</i>	<i>catch</i>	\neg <i>catch</i>
<i>cavity</i>	.108	.012	.072	.008
\neg <i>cavity</i>	.016	.064	.144	.576

For any proposition ϕ , sum the atomic events where it is true:

$$P(\phi) = \sum_{\omega: \omega \models \phi} P(\omega)$$

INFERENCE BY ENUMERATION

Start with the joint distribution:

	<i>toothache</i>		\neg <i>toothache</i>	
	<i>catch</i>	\neg <i>catch</i>	<i>catch</i>	\neg <i>catch</i>
<i>cavity</i>	.108	.012	.072	.008
\neg <i>cavity</i>	.016	.064	.144	.576

For any proposition ϕ , sum the atomic events where it is true:

$$P(\phi) = \sum_{\omega: \omega \models \phi} P(\omega)$$

$$P(\text{toothache}) = 0.108 + 0.012 + 0.016 + 0.064 = 0.2$$

INFERENCE BY ENUMERATION

Start with the joint distribution:

	<i>toothache</i>		\neg <i>toothache</i>	
	<i>catch</i>	\neg <i>catch</i>	<i>catch</i>	\neg <i>catch</i>
<i>cavity</i>	.108	.012	.072	.008
\neg <i>cavity</i>	.016	.064	.144	.576

For any proposition ϕ , sum the atomic events where it is true:

$$P(\phi) = \sum_{\omega: \omega \models \phi} P(\omega)$$

$$P(\text{cavity} \vee \text{toothache}) = 0.108 + 0.012 + 0.072 + 0.008 + 0.016 + 0.064 = 0.28$$

INFERENCE BY ENUMERATION

Start with the joint distribution:

	<i>toothache</i>		\neg <i>toothache</i>	
	<i>catch</i>	\neg <i>catch</i>	<i>catch</i>	\neg <i>catch</i>
<i>cavity</i>	.108	.012	.072	.008
\neg <i>cavity</i>	.016	.064	.144	.576


Can also compute conditional probabilities:

$$\begin{aligned} P(\neg \text{cavity} | \text{toothache}) &= \frac{P(\neg \text{cavity} \wedge \text{toothache})}{P(\text{toothache})} \\ &= \frac{0.016 + 0.064}{0.108 + 0.012 + 0.016 + 0.064} = 0.4 \end{aligned}$$

BAYES' RULE

$$P(B|A) = P(A|B) P(B) / P(A)$$

Entire Sample Space: $P(S)=1$



Area = Probability of Event

DERIVATION OF BAYES' RULE

- **Start from Product Rule:**

- $P(a, b) = P(a | b) P(b) = P(b | a) P(a)$

- **Isolate Equality on Right Side:**

- $P(a | b) P(b) = P(b | a) P(a)$

- **Divide through by $P(b)$:**

- $P(a | b) = P(b | a) P(a) / P(b)$ <-- **Bayes' Rule**

BAYES' RULE

Product rule $P(a \wedge b) = P(a|b)P(b) = P(b|a)P(a)$

$$\Rightarrow \text{Bayes' rule } P(a|b) = \frac{P(b|a)P(a)}{P(b)}$$

or in distribution form

$$P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} = \alpha P(X|Y)P(Y)$$

Useful for assessing diagnostic probability from causal probability:

$$P(Cause|Effect) = \frac{P(Effect|Cause)P(Cause)}{P(Effect)}$$

E.g., let M be meningitis, S be stiff neck:

$$P(m|s) = \frac{P(s|m)P(m)}{P(s)} = \frac{0.8 \times 0.0001}{0.1} = 0.0008$$

Note: posterior probability of meningitis still very small!