



ARTIFICIAL INTELLIGENCE



RECAP

Problem Solving Agent

- Uninformed Search
- Informed Search
- Local Search

All search schemes we covered so far were designed for single agent

- decisions of the agents determine the outcome

ADVERSARIAL SEARCH

Multiagent Environment: **Competitive**, Cooperative

Competitive environments -> adversarial search problems also called games.

Assumption(s):

- Environments is fully observable and deterministic.
- Two agents act alternately in which the utility values at the end of the game are always equal and opposite e.g., chess.

SEARCH VERSUS GAMES

- **Search:** no adversary
 - Solution is (heuristic) method for finding goal
 - Evaluation function: estimate cost from start to goal through a given node
 - Examples: path planning, scheduling activities, ...
- **Games:** adversary
 - Solution is a strategy
 - Specifies move for every possible opponent reply
 - Time limits force an approximate solution
 - Evaluation function: evaluate “goodness” of game position
 - Examples: chess, checkers, Othello, backgammon

GAME FORMULATION

Consider two players, MAX and MIN

S_0 : Initial state of the game

Player(s): Which player has to move in a state, s


Action(s): Returns the set of legal moves in a state, s

RESULT(s, a): The **transition model**, which defines the result of a move, i.e., the resulting state when action a is applied on a state s.

TERMINAL-TEST(s): A **terminal test**, which is true when the game is over and false otherwise.

UTILITY(s, p): A **utility function** (objective function/payoff function), defines the final numeric value for a game that ends in terminal state s for a player p. In chess, the outcome is a win, loss, or draw, with values +1, 0, or 0.5 .

GAME TREE (2-PLAYER, DETERMINISTIC, TURNS)



GAMES AS SEARCH


- Two players, “MAX” and “MIN”
 - MAX moves first, & take turns until game is over
 - Winner gets reward, loser gets penalty
 - “Zero sum”: sum of reward and penalty is constant
 - MAX uses search tree to determine “best” next move
- Formal definition as a search problem:
 - **Initial state**: set-up defined by rules, e.g., initial board for chess
 - **Player(s)**: which player has the move in state s
 - **Actions(s)**: set of legal moves in a state
 - **Results(s,a)**: transition model defines result of a move
 - **Terminal-Test(s)**: true if the game is finished; false otherwise
 - **Utility(s,p)**: the numerical value of terminal state s for player p
 - E.g., win (+1), lose (-1), and draw (0) in tic-tac-toe
 - E.g., win (+1), lose (0), and draw (1/2) in chess

MIN-MAX: AN OPTIMAL PROCEDURE

- Designed to find the optimal strategy & best move for MAX:
 1. Generate the whole game tree to leaves
 2. Apply utility (payoff) function to leaves
 3. Back-up values from leaves toward the root:
 - a Max node computes the max of its child values
 - a Min node computes the min of its child values
 4. At root: choose move leading to the child of highest value

OPTIMAL SOLUTION IN GAMES?

We will devise a strategy from Max's perspective, assuming that Min will play optimally!




OPTIMAL SOLUTION IN GAMES?

MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

MAX


MIN




OPTIMAL SOLUTION IN GAMES?

MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$



TWO-PLY GAME TREE



Minimax maximizes the utility of the worst-case outcome for MAX

RECURSIVE MIN-MAX SEARCH

`minMaxSearch(state)`

`return argmax([minValue(apply(state,a)) for each action a])`

Simple stub to call recursion fns

`maxValue(state)`

`if (terminal(state)) return utility(state);`

`v = -infty`

`for each action a:`

`v = max(v, minValue(apply(state,a)))`

`return v`

If recursion limit reached, eval position

Otherwise, find our best child:

`minValue(state)`

`if (terminal(state)) return utility(state);`

`v = infy`

`for each action a:`

`v = min(v, maxValue(apply(state,a)))`

`return v`

If recursion limit reached, eval position

Otherwise, find the worst child:

GAMES ARE HARD TO SOLVE!

Chess has average branching factor $b = 35$

A single game requires an average of 50 moves per player leading to 35^{100} states.

Need to do some action even when calculating the optimal decision is infeasible!

OPTIMAL MOVE?

How to choose a good move when time is limited?


Pruning allows us to ignore portions of the search tree that make no difference to the final choice

Heuristic **evaluation functions** allow us to approximate the true utility of a state without doing a complete search

ALPHA-BETA PRUNING


- Exploit the “fact” of an adversary
- If a position is **provably bad**
 - It’s no use searching to find out just how bad
- If the adversary **can force a bad position**
 - It’s no use searching to find the good positions the adversary won’t let you achieve
- Bad = **not better** than we can get elsewhere

ALPHA BETA PRUNING




Do we need to expand all nodes?

$$\begin{aligned} \text{minimax}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \\ &= 3 \end{aligned}$$



ALPHA BETA PRUNING



EXERCISE

Alpha-Beta Example

Alpha = best already explored option
along path to the root for maximizer
Beta = best already explored option along
path to the root for minimizer

